

Code Generation: x64 Assembly


(CSC 3F002 EP) Compilers

Pierre-Yves Strub


(Slide deck author: **Kaustuv Chaudhuri**)

2025-09-22 // Week 04

What You've Written So Far



Today's Focus



Agenda for Today

- A brief introduction to **x64**
 - Temporaries and the stack
 - Arithmetic and Bitwise instructions
 - Comparisons and Jumping
- Going from TAC to x64 – lab 3, week 2

Brief Introduction to x64

Assembly Program: Structure

- Three main kinds of **sections**, all *fixed size*:
 - **Text**: contains *non-modifiable* instructions
(aka: code)
 - **Data**: contains *modifiable, initialized* data
 - **BSS**: contains *modifiable, uninitialized* (0-filled) data
("Block Started by Symbol", "Better Save Space")
 - **Read-Only Data** (*rodata*): non-modifiable initialized data, usually considered part of the text section
- In an assembly file these sections can be interleaved in an arbitrary order, but assemblers usually de-interleave them
- Every section declares a list of **labels** of two varieties
 - **Local**: not exported by object code
 - **Global**: exported by object code
- In text segments, labels are **jump addresses**
- In data or bss segments, labels are **pointers**

Assembly Program: Structure Example

```
# Line comments begin with '#'
/* Can also use C/C++-style block comments */

    ## declare all the global symbols

    .globl main          # note: not .glob*a*l
    .text
main:           # note: globally accessible function
    ## instructions for main function

fib:            # note: known only within this file because no .globl
    ## instructions for fib function

    .section .rodata
msg:           # note: known only within this file because no .globl
    ## e.g., declaration of a string (i.e., byte array) named 'msg'

    .globl values
    .section .data
values:         # note: globally accessible
    ## e.g., declaration of an array named 'values'
```

A Full(y Trivial) Example

(File ret42.s)

```
## A program that only returns code 42
.globl main
.text
main:
    movq $42, %rax
    retq
```

Disassembly

```
$ as --64 ret42.s -o ret42.o          # assemble into object file
$ objdump -d ret42.o                  # disassemble object file

ret42.o:    file format elf64-x86-64

Disassembly of section .text:
0000000000000000 <main>:
 0:   48 c7 c0 2a 00 00 00    mov    $0x2a,%rax
 7:   c3                      retq

$ objdump --disassemble=main ret42.o    # can also pick function to disassemble
... same output as above ...
```

Stepping Through Assembly

- Compile the code with gcc using the `-g` flag

```
$ gcc -g -o ret42.exe ret42.s      # one shot compilation
```

- Use gdb, the system debugger

```
$ gdb ret42.exe
... bunch of output ...
Reading symbols from ret42.exe...
(gdb) l                                         shows the assembler source
1     .globl main
2     .text
3     main:
4         movq $42, %rax
5         retq
(gdb) b 5                                       set breakpoint on line 4
Breakpoint 1 at 0x7: file ret42.s, line 5.
(gdb) r                                         begin execution
Starting program: /home/kaustuv/....ret42.exe

Breakpoint 1, main () at ret42.s:5
5         retq
(gdb) i r rax                                    display contents of RAX
rax          0xa                               42
(gdb) c                                         continue
Continuing.
[Inferior 1 (process 105074) exited with code 052]
(gdb)
```

Computer Memory: Quick Recap

- Memory: stores data, addressable
 - Volatile memory: contents invalidated on power loss
 - Persistent memory: contents preserved on power loss (aka *non-volatile* memory)
- A computer has many kinds of memory
 - Storage: hard disks, flash memory, etc. (persistent)
 - Seek latency: spinning disk: 3ms, SSD: $16\mu\text{s}$
 - Main Memory: memory modules, variable (volatile)
 - Reference: 100ns
 - Cache: fixed, part of the processor die (volatile)
 - L2 reference: 4ns, L1 reference: 0.5ns
 - Registers: fixed, hard-wired in the processor logic (volatile)
 - Latency: depends on clock speed. E.g., 4GHz proc. has $\leq 0.25\text{ns}$
 - Fully parallelized: every bit accessed in parallel

link: [excellent interactive visualization tool for latency](#)

x64 Registers

- General purpose registers (GPR) – stores integers

RAX RBX RCX RDX RSI RBP RSP RDI
R8 R9 R10 R11 R12 R13 R14 R15

- Floating point registers (FPR) – stores floats

FPR0 FPR1 FPR2 ... FPR7

- Special purpose registers (SPR)

RIP RFLAGS

- MMX registers (64 bits) – SSE

MMX0 MMX1 MMX2 ... MMX7

- XMM registers (128 bits) – SSE2

XMM0 XMM1 XMM2 ... XMM7

- YMM registers (256 bits) – SSE3, SSSE3

YMM0 YMM1 YMM2 ... YMM7

- ZMM registers (512 bits; not yet consumer-grade) – AVX-512

ZMM0 ZMM1 ZMM2 ... ZMM7

GPRs and the C Calling Convention (ABI)

- Some of the GPRs have special interpretations
- **Stack frame:**
 - RBP: *base pointer*, points to *end* of allocated stack and start of book-keeping data
 - RSP: *stack pointer*, points to *start* of allocated stack
 - Don't use these two registers for computations!
- **Callee-save registers:** (“*non-volatile*”)
 - Must retain values across function calls
 - In the C (SysV) ABI: RBX, RBP, R12, R13, R14, R15
 - Don't use any of these for lab 3 – we'll come back to them later
- **Caller-save register:** (“*volatile*”)
 - Any GPR that is not callee-save
 - Callers must keep copies if the old values matter to them
 - Called procedures can overwrite them...
... so, feel free to use them for lab 3

Move Immediate

```
movq    $42, %rax  
immediate destination
```

- Stores a given number (“*immediate*”, prefixed \$) in a register
- Note: the immediate must fit in 32 bits!
- Immediate value can be given in Hexadecimal (e.g., \$0x42)

Move (Copy) Register

```
movq    %rax, %rdx  
        source destination
```

- Both source and destination register must be the **same size**
- If source and destination are the same, this is a null operation

Addition and Subtraction (Registers)

```
addq    %rdx, %rax  
        source destination
```

- Adds the value of the source register to the value of the destination register
(think: $\%rax = \%rax + \%rdx$)
- Overflow flags are set based on destination $(\sim \text{lab 3, checkpoint 2})$
- Subtraction (**subq**) defined similarly.

Addition and Subtraction (Immediates)

```
addq    $42,    %rax  
immediate destination
```

- Adds the immediate to the value of the destination register (think:
 $\%rax = \%rax + 42$)
- Subtraction (`subq`) defined similarly.

Arithmetic Negation

```
negq      %rax  
          source/dest
```

- Negates the value of the source register and stores it back in the same register (think: $\%rax = - \%rax$)

Signed Multiplication

```
imulq    $42,    %rax  
          factor   destination
```

- Multiplies the value of the factor to the destination register.
- There is a more complicated single operand version of `imulq` that has much higher precision, but you may choose to ignore it for this course.

Signed Division and Modulus

<code>idivq</code>	<code>%rcx</code>
	divisor

(3 bytes)

- Divides the value of `RDX:RAX` by the value of the divisor register, and stores the quotient in `RAX` and remainder (modulus) in `RDX`
- Think:

```
%rax = (%rdx:%rax) / %rcx  
%rdx = (%rdx:%rax) % %rcx
```

- Because this can obliterate `RDX`, you may want to save and restore its value using a book-keeping register (e.g., `R11`)
- To sign-extend a value stored in `RAX` to `RDX:RAX`, use the nullary instruction: `cqto`.

Shifting Left

(shifting right is similar)

salq	\$7,	%rax
	count	destination
salq	%cl,	%rax
	count	destination

- Shifts the destination register left by the count specified as an immediate (which must fit in 6 bits) or in the register **CL** (which is the low 8 bits of **RCX**)
(Think: $\%rax = \%rax \ll 7$)
- Note: no other register besides **CL** can be used as the count

Bitwise And, Or, Xor

```
andq    $42,    %rax  
        imm. src. destination
```

```
andq    %rdx,    %rax  
        reg. src. destination
```

- Computes the bitwise and of the source (either an immediate or the value of a register) and the value of the destination register, and stores the result in the destination register. (Think: `%rax = %rax & %rdx`, etc.)
- Similarly for or (`orq`) and xor (`xorq`)

Bitwise Complement

```
notq      %rax  
          source/dest
```

- Takes the bitwise complement (aka 1's complement) of the value of the source register and stores it back in the same register
(Think: `%rax = ~ %rax`)

Other Instructions

- The collection of instructions presented so far will suffice for all **arithmetic** and **bitwise** operations with registers and immediates in this course
- However, you may wish to use some additional instructions that are optimized for specific use cases. Some suggestions:
 - `inc` and `dec`: increment and decrement by 1
 - `movsx` and `movsxd`: sign extend to/from arbitrary registers, not just `RAX`
 - `movzx`: move with zero-extension
 - `andn`: “and not” (Haswell microarchitecture and later)
 - `rol`, `ror`: rotate left/right

Labels and Instruction Pointer

- **Pointers:** (virtual) memory addresses as values
Can be dereferenced to read/write memory
- **Label:** addresses of an instruction in compiled code
- **Instruction pointer (RIP):** contains the address of the *next* instruction, so setting **RIP** changes the code path
- **Jump:** specialized instructions to set **RIP**.
 - **jmp:** absolute jump
 - **jz, jnz, jl, jle,** etc: conditional jumps

Comparison Instruction

```
cmpq      $42, %rax  
          src2    src1  
          (reg/imm) (reg/mem)
```

- Sets the status flags register (`%eflags`) based on $\text{src1} - \text{src2}$
- Note: does not affect any of the other registers
- Overlays:
 - `cmpl` $\$42, \%eax$
 - `cmpw` $\$42, \%ax$
 - `cmpb` $\$42, \%al$

Jump Instructions

```
jmp    .L42  
jcnd   .L42  
label
```

- Jumps unconditionally (`jmp`) or when a particular condition flag is set/unset based on an earlier `cmpq`
- Some conditional jumps: (note synonyms)


Instruction	Condition	wrt: [<code>cmpq</code> <code>src2, src1</code>]
<code>je, jz</code>	$ZF=1$	$src1 - src2 == 0$
<code>jne, jnz</code>	$ZF=0$	$src1 - src2 != 0$
<code>jle, jng</code>	$ZF=1 \text{ or } SF \neq OF$	$src1 - src2 \leq 0$
<code>jl, jnge</code>	$ZF=0 \text{ and } SF \neq OF$	$src1 - src2 < 0$
<code>jge, jnl</code>	$ZF=1 \text{ or } SF \neq OF$	$src1 - src2 \geq 0$
<code>jg, jnle</code>	$ZF=0 \text{ and } SF=OF$	$src1 - src2 > 0$

Storing Temporaries in the Stack

- There are a finite number (≤ 14) of usable registers.
What if we have more variables/intermediates/temporaries?
- Answer: we *spill* some temporaries onto the stack
- General spilling algorithm:
 - ① Assign one register (say R11) for book-keeping spilled temps
 - ② Allocate a *slot* on the stack for each temporary that is to be spilled.
 - ③ Whenever that temporary is needed, load it from the stack slot into R11, perform the operation, then save R11 back into the stack slot.
- Sometimes the stack slot can itself be an operand (stay tuned!)

The Stack Frame, Briefly

(we'll cover the stack in full gory detail in lab 4)



Using the Stack

- Until lab 4 (procedures), we will ignore:
 - Stack parameters
 - Return address
- All we need for now: there are *n* stack slots between RSP and RBP available to be used for temporaries.
- What is *n*?
 - First approximation: number of temporaries in TAC
 - Better: temporaries can be allocated to caller-save registers that are not used for anything else (yet): RSI, RDI, R8 – R11.
 - Warning! RBX, R12 – R15 are callee-save!

Allocating the Stack

(Hint: use this as a template for lab 3)

```
1      .globl main
2      .text
3 main:
4      # store old RBP at top of the stack
5      pushq %rbp
6      # make RBP point to just after (i.e., top) stack slots
7      movq %rsp, %rbp
8
9      # Now we allocate stack slots in units of **16** bytes (= 128 bits)
10     # E.g., for 7 slots, we would allocated 8 * 8 = 64 bytes
11     subq $64, %rsp
12
13     #
14     # The rest of the compiled code from BX1
15     #
16
17     # restore RSP and RBP
18     movq %rbp, %rsp
19     popq %rbp
20
21     movq $0, %rax      # return code 0
22     retq
```

Loading and Storing to Stack Slots

syntax for stack slots as operands

Stack slot 1: -8(%rbp)

Stack slot 2: -16(%rbp)

Stack slot 3: -24(%rbp)

⋮

Stack slot n : - $8n$ (%rbp)

`movq -8(%rbp), %rax` (load slot into register)
source destination

`movq %rdx, -8(%rbp)` (save register to slot)
source destination

- Can use a dedicated register (e.g., `R11`) for load/store to stack slots followed/preceded by a register-register `movq`.
- Both source and destination cannot be stack slots!

Optimization: Using Stack Slots Directly as Operands

```
addq    -8(%rbp), %rax  
        source      destination  
  
addq    %rdx, -8(%rbp)  
        source      destination
```

- A stack slot be used instead of a register operand when that register is **arbitrary**
- Note: not the case for **sarq**, **idivq**, etc. which will *always* use specific registers
- All operands cannot be stack slots!

Lab 3, Checkpoint 1

Lab 3 (control structures), Checkpoint 1

- Lab 3 is about compiling BX control instructions to x64, ...
 - Checkpoint 2: converting TAC (+labels, jumps) to x64
 - Note: In x64, destination reg = final arg reg (usually)
Translating TAC to x64 is not completely trivial
 - May need one or two book-keeping registers (use R11 unless you need something else)
 - Reminder: Lab 3 has a full duration of **2 weeks**
-