

[CSE301 / Lecture 0]

An introduction to functional programming

Noam Zeilberger

Ecole Polytechnique

3 September 2025

What is functional programming?

Hard to give a precise definition, but a rough approximation is that functional programming is a *style of programming that emphasizes function application and function composition.*

Function application

“The act of evaluating a function on some argument.”

E.g., apply $f(x) = \sqrt{x}$ to 2, obtaining $f(2) \approx 1.41421$.

Depending on the compiler, application may be implemented using lower-level operations like pushing arguments onto a stack, etc. – although sometimes we need to know these details, usually we can treat function application as a higher-level abstraction.

Function composition

“The act of combining two or more fns. to define a new function.”

E.g., given $f(x) = \sqrt{x}$, $g(x) = \sin x$, $h(x) = e^x$, define

$$i(x) = h(f(x) + g(x))$$

Observe we can compose both “in sequence” and “in parallel”.

Again, a compiler may need to make additional choices (e.g., store $f(x)$ before computing $g(x)$? vice versa? multicore?), but the functional notation nicely captures just the logical dependencies.

What makes a programming language “functional”?

Possible to program in the functional style in almost any language, but a *functional programming language* makes it easier.

Typically, by including at least some of the following features:

pattern-matching higher-order functions rigorous typing

(This is not an “official” list. But you will hopefully come to appreciate why these three features are especially useful.)

A few examples of languages with all these features:

Haskell, OCaml, Rocq, Agda, Rust, Lean, ...

(Quick poll: who has used any of these languages?)

Why learn functional programming?

FP had a reputation as “academic” for a long time, but Haskell and OCaml have been used in industry for at least two decades, and FP concepts are increasingly going mainstream.

Some practical benefits of FP:

- Powerful notations inspired by mathematics and logic
- Better control over “side-effects” of functions
- In principle, easier to parallelize

Overall, FP simplifies the task of going from an abstract description of a problem to an efficient and reliable implementation in code.

... But also: *it's beautiful!*

An example

$\text{part} :: (a \rightarrow \text{Bool}) \rightarrow [a] \rightarrow ([a], [a])$

$\text{part } p [] = ([], [])$

$\text{part } p (x : xs) = \text{if } p x \text{ then } (x : ts, fs) \text{ else } (ts, x : fs)$

where

$(ts, fs) = \text{part } p xs$

$\text{qsort} :: \text{Ord } a \Rightarrow [a] \rightarrow [a]$

$\text{qsort } [] = []$

$\text{qsort } (x : xs) = \text{qsort } left ++ [x] ++ \text{qsort } right$

where

$(left, right) = \text{part } (\lambda y \rightarrow y < x) xs$

A brief (pre-)history

1920s-30s: Alonzo Church and his students Kleene and Rosser develop λ -calculus.

1937: Alan Turing proves equivalence between TM-computability and λ -definability.

late 1950s: John McCarthy develops the LISP language.

mid 1960s: Peter Landin promotes λ -calculus as a conceptual tool for reasoning about programming languages.

1970s: striking connections between programming, logic, & math!

Unwinding the Curry-Howard-Lambek correspondence

In the tumultuous 1970s (and late '60s):

- Dana Scott invents *domain theory*
- Jean-Yves Girard & John Reynolds both independently discover the *polymorphic λ -calculus*
- J. Roger Hindley and Robin Milner both independently discover an algorithm for *polymorphic type inference*
- Per Martin-Löf introduces *dependent type theory*
- Joachim Lambek's work on *cartesian closed categories*, building on Bill Lawvere's earlier work on categorical logic, as well as Lambek's own older work in mathematical linguistics

More recent history

also in the 1970s: Guy Steele & Gerald Sussman develop Scheme and write “Lambda: The Ultimate” series of papers

also also in the 1970s: Robin Milner and others develop ML

1980s: ML evolves into Standard ML and Caml (later OCaml)

1987: an international committee starts work on Haskell

1989: first release of the Coq/Rocq proof assistant

1992: Phil Wadler’s “Monads for functional programming”

1996: OCaml developed by Xavier Leroy, Jérôme Vouillon, et cie

2007: first release of Agda proof assistant, written in Haskell

2021: “LAMBDA: The ultimate Excel worksheet function”

Why Haskell for this course?

An elegant language with a rich ecosystem. (So is OCaml.)

Haskell is a **pure & lazy** functional programming language:

- Purity forces you to think more rigorously about side-effects.
(Though question of *how* to think about them is still open...)
- In retrospect, laziness was probably a bad idea, but at least it is an interesting one! (We will study it, but not emphasize it.)

Ultimately, we will use Haskell as an *intellectual tool* for learning about functional programming, although one day maybe you will find it practically useful!

Course practicalities

Use Moodle for:

- Handing in assignments
- Receiving announcements

Use Slack for:

- Questions about lectures / labs / projects
- Technical support questions

The course webpage (<https://noamz.org/teaching/CSE301/>) has:

- Practical information about assessment etc
- A provisional schedule
- Lecture notes and slides
- Lab descriptions